

The Athomaris Library

A Framework for Business Processes, Database Integration, and Web Interfaces

Thomas Schöbel-Theuer

Version 0.3

Features:

- Automated **business processes: orchestration level**, and execution by a very generic **state transition machine**, with generic interfaces to almost arbitrary service providers.
- Provides **temporal databases**: the **history** of all changes may be recorded, even deletions. Temporal queries possible, in addition to non-temporal views of the data.
- **Versioned schema management**: when creating new versions of **\$SCHEMA**, SQL statements like **alter table** are generated. Live data is left intact whenever possible.
- **Access management**: fine-grained access control based on user profiles. The user interface obeys these rules by presenting only those fields and buttons for which access is currently allowed.
- **Distributed access** to heterogeneous databases. Logical integration of multiple data sources into a single “**virtual database**” with **cross-join capabilities** and **cross-referential integrity** are in preparation.
- **Scalable**: transparent support for master/slave replication setups in large enterprises. **Vertical splitting** via views. Transparent **horizontal splitting** is in preparation.
- **Automatic consistency updates**: tables on different (distributed) databases may be kept always in-sync with each other, either symmetrically or asymmetrically in multiple modes.
- Generic **default user interface**: once **\$SCHEMA** has been created, the database can be accessed instantly via a customizable web-based interface (without programming).
- Internal **Model-View-Controller** (MVC) architecture: generic **template management** with internationalization and user-profile support. Customizable and extensible.
- **Security**: generated SQL code and templates are automatically escaped to prevent attacks like SQL injection, XSS, etc.

Contents

1	Basic Setup	3
1.1	Demo Projects	3
1.2	New Projects	5
2	Basic configuration	6
2.1	Basic \$SCHEMA configuration	6
2.1.1	Table Definitions in \$SCHEMA	6
2.1.2	FIELDS	7
2.1.3	Temporal tables	9
2.1.4	Initializing Data	10
2.1.5	Views	10
2.2	\$CONFIG configuration	10
3	Basic Customization	12
3.1	Link-Headers	12
3.2	Generic Application Layer	12
3.3	Generic Template Mechanism	13
3.4	Default Templates	15
4	Business Process Engine	16
4.1	Concepts	16
4.2	Usage	16
4.3	Tables for the Orchestration Level	17
4.4	Pseudo Events	19
5	Advanced Features	20
5.1	Automatic Data Synchronization	20
5.2	Subrecords and Display/Editing of References	20
5.3	Programmers API to the Database	20
5.3.1	Data Format	20
5.3.2	Update Operations	20
5.3.3	Reading of Data	21
5.3.4	Full Queries / Subqueries	22
5.3.5	Aggregated Queries / Subqueries	22

1 Basic Setup

Requirements: you need Linux, Apache with `mod_php5` and some of the supported databases (currently only MySQL, but a Sybase driver and other drivers are planned). If you want to use the business process execution engine, you also need to install and enable the `pcntl` extension module of PHP.

Copy the sources of the Athomaris PHP Library to the document root of your webserver. In the following, we will assume that this is `/www`. If you use another place, you have to adapt the following examples accordingly. So it looks like

- `cd /www; sudo tar xzpf athomaris.tar.gz`

This creates a directory `/www/athomaris/`. It is important to do this as `root`, because there are directories owned by different users¹. Afterwards, configure Apache the following way:

```
<Directory "/www/athomaris">
    Options ExecCGI FollowSymLinks
    AllowOverride AuthConfig
    AuthType Basic
    AuthName "The Athomaris PHP Library"
    AuthBasicProvider file anon
    AuthUserFile /dev/null
    Require valid-user
    AuthBasicAuthoritative Off
    Anonymous_NoUserID off
    Anonymous *
    Order allow,deny
    Allow from all
</Directory>
Alias /demo_basic "/www/athomaris/demo_basic"
Alias /demo_advanced "/www/athomaris/demo_advanced"
Alias /demo_business "/www/athomaris/demo_business"
```

Make sure the `authn_anon` module of Apache is loaded. Many Linux distributions have it disabled by default. Typically, you need to add it to `/etc/sysconfig/apache2` or to some file in `/etc/apache2/sysconfig.d/` or to another place within `/etc/apache2/` (distribution specific). Afterwards don't forget to restart Apache!

Important: if you communicate over **untrusted networks**, it is *highly recommended* to use **https** instead of **http**. Otherwise an attacker will be able to watch your passwords from the basic HTTP authentication. You will need an SSL certificate for **https**. Creating or obtaining certificates and configuring Apache for **https** is beyond the scope of this paper. Please consult the Apache documentation and various web resources for that.

Further potential pitfalls: if you want to process large data fields such as TEXT or BLOBs, you probably have to increase the *memory limits* of PHP. Some distributions use rather small defaults.

1.1 Demo Projects

Provided you have a local MySQL installation running, you can immediately use the demo projects `/www/athomaris/demo_basic/`, `/www/athomaris/demo_advanced/`, and `/www/athomaris/demo_business/`.

¹If it went wrong, or if you got the sources via `svn`, you can fix the permissions as follows:

- `cd /www/athomaris; chmod 755 */compiled; sudo chown -R wwwrun */compiled/`

Inspect the files in the demo project. There are not many, and they are rather small and simple. For example, `config.php` in each of the demo directories describes the connections to database servers and their drivers. The others mostly include generic code from the `/www/athomaris/common/` subdirectory.

The database schema in `/www/athomaris/demo_basic/schema/schema01.php` is the most central place you will have to deal with. Its content should be rather self-explanatory for PHP programmers. Details on it will be provided later.

Use a web browser such as Firefox or Konqueror to open the following URL:

- http://localhost/demo_basic/create_schema.php

If your Apache setup is correct, your browser will ask you for a username and a password (using the basic authentication mechanism of the HTTP protocol). Type in `root` and the appropriate root password for your MySQL server. Note that you *must* have already set a root password for MySQL (which is not the same as the root password of the operating system), and you normally really *need*² root access in order to be able to create a new database³. If you have trouble with that, consult the MySQL and the Apache documentation, inspect the logfiles, and probably turn on debugging of some of the components.

After successfully gaining root access, you will see a *preview* screen with MySQL commands. Look at them. If you want to actually execute these commands, click on the button named `demo_advanced`. If all is ok, you will be notified that the database has been created.

If you like, you may check the database by hand, for example:

- `mysql -p -u root demo_basic`
- `show tables;`
- `describe foos;`

You can see that the schema definition from `schema/schema01.php` has been applied pretty straightforward. However, some additional columns have been *automatically* created. Some of them have to do with *temporal databases* (see section 2.1.3).

Now let us change the schema, by adding a new column. Copy the file `schema/schema01.php` to `schema/schema02.php` without altering the old version `schema01.php`, and afterwards edit the new version `schema02.php` in the following way:

Copy the definition of `foo_name` including its sub-structure, and rename it to `foo_somethingelse`. You may change the `TYPE` field to a different type, such as `varchar(50)`, or change the `DEFAULT` value to something you like. A list of available types will be provided later. Note that you must always provide a `DEFAULT` value when you want to create a *new* column in the database. You can also change the `DEFAULT` of an existing column, or even remove a column (but please don't remove columns mentioned in `UNIQUE` keys - they are vital).

After modification of `schema02.php`, just open the URL for `create_schema.php` again. Now you will see different SQL statements, in particular `alter table` or `drop column` statements, depending on what you have changed between `schema01.php` and `schema02.php`. When executed, these statements will update the database to the new version of the schema.

What's the "big clue" with that? When using advanced features such as `profiles` (see `demo_advanced`) or when implicitly adding further tables (e.g. for the business process engine, see section 4), these tables will be automatically *kept consistent* with the global schema \implies aka **schema management**. In future Athomaris releases, we want to provide web interfaces for schema management, then you no longer will need to deal with PHP variables.

In order to use the database, open the following URL:

- http://localhost/demo_basic/

²Exception: you are a database expert who knows how to grant appropriate privileges to other users.

³Some Linux distros deliver MySQL without any root password. In such a case Athomaris will not work, because passwords are mandatory. You may use the following command for setting an initial password:

- `mysqladmin -u root password "secret"`

You can enter new tuples into the database, browse the data in various orders and selections, update and delete the data as you like. It should be rather self-explanatory what you can do. Just try it!

If you like, you may change the schema again using the name `schema03.php`, after you already have populated your tables with data.

If you want to customize the user interface, look at the file `/www/athomaris/common/tpl/generic/generic.tpl` and copy some of its template definitions over to `/www/athomaris/demo_basic/lang/generic/generic.tpl`. There you may modify the HTML code (details on the template syntax, macro substitutions, hooks are in section 3.3). After having done that, you must restart the template compiler via the following URL:

- `http://localhost/demo_basic/translate_tpl.php`

When run successfully, it translates all the templates to PHP code which you can find in `compiled/generic_generic.php`. This code does not aim to be readable, but is very fast on execution.

1.2 New Projects

The easiest way to start a new project is by copying one of the demo projects (or any other project) to a new subdirectory. For example

- `cd /www/athomaris/; sudo cp -a demo_business myproject`

It is important to use the `-a` flag as `root`, in order to keep some symlinks intact and to ensure that the subdirectory `compiled/` remains owned by the Apache user and is writeable by it.

Afterwards, you should configure Apache such that your new directory is accessible via HTTP over some URL. For example, add the following to `httpd.conf`:

- `Alias /myproject "/www/athomaris/myproject"`

Don't forget to restart Apache after that! You should be able to access the new project over the following URLs:

- `http://localhost/myproject/create_schema.php`
- `http://localhost/myproject/translate_tpl.php`
- `http://localhost/myproject/`

As always with PHP projects, you may add any new files and subdirectories as you like. To ease your life as a developer, most of your work can be delegated to the Athomaris PHP library just by adding the following lines:

- `require_once("config.php");`
- `require_once("$BASEDIR/../../common/app.php");`

and calling only a few functions. Details are explained later. If you want to learn how the default page layout works, inspect `common/generic.php`.

2 Basic configuration

2.1 Basic \$SCHEMA configuration

The subdirectory `schema/` contains files of the form `schemaxx.php` where `xx` must be a number. The only purpose of these files is to define different versions of `$SCHEMA` and `$EXTRA`. In this section, only `$SCHEMA` is explained.

`$SCHEMA` is a nested PHP structure.

The outmost level of `$SCHEMA` is a PHP hash indexed by *table names*, obeying standard identifier syntax. All table names in the system must be globally unique. By default, the PHP table name is also used as SQL table name. This default association can be overridden via the `SQLNAME` attribute (see later).

2.1.1 Table Definitions in \$SCHEMA

Each table is described by a nested PHP hash which may contain some of the following attributes:

`FIELDS =>` *hash* describing the database columns. No default.

This sub-structure is described in section 2.1.2.

`TEMPORAL =>` `true` or `false`. Defaults to `true` when omitted.

When a table is created as a temporal table, some special `FIELDS` are automatically added by default. Details are described in section 2.1.3.

When `TEMPORAL` is `false`, no `FIELDS` are automatically added, not even the default primary key `tablename_id`. Thus you *must* create that explicitly, and you *must* add a `PRIMARY` definition (see below).

`PRIMARY =>` *string*. Name of the primary key. Defaults to `singulartablename_id` when omitted.

Normally this attribute is determined by default. When `TEMPORAL` is `false`, or when interfacing with external databases having different schema conventions, you can set a different name here.

Hint: recommended *best practice* is to use the default only as a *hidden* primary key. This means that the *end-user* should never use its value for any purpose; he/she should not even *see* this value. Instead, the end-user should only see and use `UNIQUE` keys as described below. This has tremendous advantages: defining the user-visible identification attributes as *dependent attributes* allows to change them at runtime. When combined with `REFERENCES` and “*on update cascade*” (see below), your system becomes flexible in a way which is *vital* for large enterprises.

Exception when interfacing to external databases: if you don’t have control over the schema (e.g. when connecting to *existing* databases), you can also specify a *combined key* by enumerating the column names separated with comma.

`DB =>` *string*. References a database name from `$CONFIG`. When omitted, it defaults to the first entry in `$CONFIG`.

Use this when you maintain multiple database connections in parallel. Details on `$CONFIG` are in section 2.2.

`ENGINE =>` *string*. Only used for MySQL. Defaults to `MyISAM` when omitted.

`UNIQUE =>` *array* of strings. Names of secondary keys. No default. Can be omitted.

Tells which columns or column combinations should be uniquely indexed. If you want to create a combined index, just use a string where the column names are comma-separated.

`INDEX =>` *array* of strings. Names of indexes. No default. Can be omitted.

Works like `UNIQUE`. The only difference is that non-unique indexes are created.

ACCESS => *string* with one of the values ‘**r**’, ‘**R**’, ‘**w**’, ‘**W**’, or ‘**n**’. Defaults to ‘**W**’ when omitted.

Maximum access rights for the *whole table* which cannot be exceeded by anybody except **root**. The access right codes have the following meaning:

n no access at all.

r read access to the database, but not displayed to the user GUI.

R read access to both the database and for the user.

w write access to the database, but the user can only read it over the GUI.

W write access for both the database and for the user.

REALNAME => *string*. Defaults to the tablename when omitted.

The physical table name as used by the backend database may differ from the *logical* name used by Athomaris. By widely using this feature, you may ease *schema integration* in heterogeneous environments, such as company mergers etc. When combined with the cross-join and cross-referential integrity capabilities of Athomaris, **logical integration** of large systems is possible.

SINGULAR => *string*. When omitted, it defaults to the tablename without trailing “s”.

The **Athomaris convention** is to use English plural forms for (logical) table names. In contrast, column names should always start with the corresponding singular form whenever they *logically belong*¹ to that table. This way, (generated) SQL code becomes better understandable for humans (and foreign code bypassing the Athomaris PHP library becomes more intuitive). For example, table **foos** uses column names such **foo_id** (here as the primary key). Because there are some linguistic exceptions in English (e.g. “classes” → “class”), you can set the singular form here. The singular is used as a basis for automatic creation of column names, such as ***_id**.

However, when interfacing to external / existing databases, you often have to obey foreign conventions conflicting with the default systematics of Athomaris. In such a case, you can use the following directives for further fine-grained control (the meaning of most of them is described in section 2.1.3):

FIELDNAME_ID => *string*. Defaults to “*singular_id*”.

This is not necessarily the same as **PRIMARY**, because **PRIMARY** may be a *combined key*, while **FIELDNAME_ID** specifies which column should get an **AUTO_INCREMENT** attribute.

FIELDNAME_VERSION => *string*. Defaults to “*singular_version*”.

FIELDNAME_DELETED => *string*. Defaults to “*singular_deleted*”.

FIELDNAME_MODIFIED_FROM => *string*. Defaults to “*singular_modified_from*”.

FIELDNAME_MODIFIED_BY => *string*. Defaults to “*singular_modified_by*”.

... NYD

2.1.2 FIELDS

TYPE => *string*. No default. Cannot be omitted.

describes the SQL type of the column. Valid types are:

int, **smallint**, **bigint**, **int(*n*)** Integers (with number of decimal digits *n*).

bool, **boolean** Integers with allowed values 0 and 1.

char(*n*), **varchar(*n*)** Strings, with length or maximum length *n*.

datetime, **time**, **timestamp** See corresponding SQL types.

¹Here are some exceptions which don’t “logically belong” to some given table:

- Foreign keys (see **REFERENCES** in section 2.1.2): please use the name from the referenced table. This way, almost all possible joins can be expressed as *natural joins*, and the schema is self-documenting for humans.
- Anything which references to external information, such as filesystem data. This way, you explicitly express that the corresponding *domain* is not under your control. If you have many such cases, it is wise to even introduce some systematics for that.

text, **tinytext**, **mediumtext**, **longtext** Most databases store variable-length texts columns in a seaparate heap, not in the data record. This affects nonfunctional access properties at runtime.

blob, **tinyblob**, **mediumblob**, **longblob** Ditto. In difference to **text**, no comparion for equality is possible.

Note: although some of these types look like being **mysql**-specific, other drivers such as **sybase** will internally convert them to the most closest type available there.

DEFAULT => *string*. No default. Can be omitted, but strongy discouraged.

Best practice is to always provide defaults for any column. The default should clearly indicate that initialization has been omitted. This way, you can see when somebody has used the default.

The *string* must denote SQL code. If you want NULL as default value, just denote “**null**”. If you want a string, you have to provide the quotes within the string, such as “**‘text’**”.

CHANGE_FROM => “*old_columnname*”. No default, should be omitted if not necessary.

Specifies that a column name has changed between two schema revisions (in **schema/schemaxx.php**). In essence, this will be translated to **alter table ... change column ...** statements. Attention! Don’t forget to remove the **CHANGE_FROM** assignment in the *following* revision, otherwise it would try to rename the (then non-existing) old column once more!

REALNAME => *string*. Defaults to the fieldname when omitted.

The physical column name as used by the backend database may differ from the *logical* name used by Athomaris².

LENGTH => *array(a, b)*. When omitted, *array(0, maxlen)* is used as default where *maxlen* stems from the **TYPE** definition.

Specifies the minium and maximum allowed string lengths *a* and *b*. This is translated to a **check** clause in some SQL dialects, and to a runtime check where SQL is incapable of directly handling this.

REGEX => *string*. No default. May be omitted.

Create runtime checks (or in some SQL dialects even **check** clauses) which ensure that the value obeys the given regular expression. The regular expression must be in Perl syntax, enclosed in “*/.../*”.

REFERENCES => *hash*. No default. May be omitted.

Specifies that referential integrity between tables must be obeyed at runtime. This works even with MySQL versions which don’t support referential integrity such as MyISAM tables (appropriate checking code will be automatically executed by the Athomaris PHP Library at runtime). The *hash* must obey the following structure:

“*tablename.fieldname*” => *array(keystrings)*. The *.fieldname* may be omitted when it has the name as the current column (corresponding to a *natural join* possibility). The *keystrings* must be one or more of the following indicating the well-known ANSI SQL meaning:

- on delete cascade**
- on delete set null**
- on update cascade**
- on update set null**

ACCESS => *string*. Defaults to “**W**” when omitted.

Specifies maximum access rights for the current column. The meaning of the code letters is the same as already described at 2.1.1. The only difference is the granularity level, which is column granularity here.

OPTIONS => *string*. No default. Should be omitted.

This allows direct throughpassing of MySQL options. Try to avoid this!

²In essence, this automatically translates to SQL code like **select *physicalname* as *logicalname*, ... from ...**

2.1.3 Temporal tables

When a table `foos` is created with `TEMPORAL => true` (or by default), the following columns are automatically added to the table:

`foo_id => array("TYPE"=> "bigint", "DEFAULT" => "auto_increment").`

This is the primary key of the table. Best practice is to use this *only* as “physical object id” of data records, but **never** for *logical* identification! Logical identification should always be done via additional secondary keys (see `UNIQUE` in section 2.1.1). This has the major advantage that the logical identification can be *changed* at any time, even consistently across the whole system if you define appropriate `REFERENCES` with `on update cascade` on them (see section 2.1.2).

`foo_version => array("TYPE"=> "timestamp", "DEFAULT" => "current_timestamp").`

The timestamp of any changes to this table is automatically recorded in this column. Thus you may use it for inspection of historical versions of tuples.

`foo_deleted => array("TYPE"=> "boolean", "DEFAULT" => "false").`

Normally this value is always `false`. When a tuple is *logically* deleted, it is not deleted in reality, but rather *marked* as deleted via this column. This way, you can later inspect even deleted tuples via the temporal table `foos_tp` (see below).

`foo_modified_from => array("TYPE"=> "varchar(16)", "DEFAULT" => "null").`

This column automatically records the hostname or ip address of anybody who does any modification on some tuple, provided that the access is done via the PHP programming interface such as `db_insert()` (see section 5.3). If you bypass the default programming interface with your own hand-written SQL code, it is highly advisable to provide a value here, because this column is essential for humans if they try to comprehend anything in a complex business application.

`foo_modified_by => array("TYPE"=> "varchar(16)", "DEFAULT" => "null").`

Automatically records the login username of interactive users, or the scriptname of PHP scripts when they do any modification on some tuple.

When defining a table `foos` as temporal, there will be two tables created behind the scenes:

- `foos` is a SQL view providing the non-temporal view on the table. “Non-temporal” means that you cannot access old versions of a tuple, just as with ordinary flat SQL tables where old values are always overwritten by updates. From a user’s perspective, this has (almost) the same semantics as a usual SQL table, as if there were no temporal extensions at all. The only exception is updating: a drawback of MySQL is that it does not allow updates to most views. Therefore you must use `foos_tp` instead of `foos` for updates if you want to bypass the default programming interface with your own hand-written SQL code.
- `foos_tp` is the actual table where all the temporal history is stored. When a tuple is *logically updated* (e.g. via `db_update()` c.f. section 5.3), a *new* tuple is *inserted* into `foos_tp` behind the scenes instead (aka COW = Copy On Update strategy). The new tuple has the same `foo_id`, but a newer `foo_version` timestamp.

The relation between `foos_tp` and `foos` is illustrated by the following SQL code:

```
create view foos as
select * from foos_tp t1 where t1.foo_version in
(select max(t2.foo_version) from foos_tp t2
 where t1.foo_id = t2.foo_id and not t1.foo_deleted
);
```

As you can see, this filters out any old versions, as well as any tuples which have been *logically* deleted. However the “deleted” tuples remain accessible via `foos_tp`.

2.1.4 Initializing Data

In order to populate your database with some initial data, you may use the global variable `$INITDATA` (see examples in `demo_basic` and `demo_advanced`). This is handy for startup tasks such as creating initial profiles and users. If you omit initial data for tables `profiles`, `languages`, or `users`, some resonables defaults will be supplied (try `demo_advanced` for details).

`$INITDATA` is a hash, indexed by table names. Each table is associated an array of records. In turn, each record is a hash which associates field names to field values.

Note that initializing is only done on `create_schema.php` (see section 1.1) when there is exactly one `schema/schemaxx.php`. When updating the schema to a newer version, you have to update your data by hand (this is necessary because the initial data may be completely altered in the meantime).

2.1.5 Views

Athomaris can create and manage SQL views in `$SCHEMA`. Views may be built upon arbitrary queries, which may join multiple tables as described in sections 5.3.4 and 5.3.5. A View may be used like any other table, but with some restrictions:

1. Views must not span tables from different `$CONFIG` hosts (see section 2.2).
2. Views are only readable (equivalent to `ACCESS => 'R'`).
3. There is no primary key (equivalent to `PRIMARY => null`). As a consequence, no `REFERENCES` to views are possible.
4. A view cannot reference to the temporal version of another table, but only to the flat one. Instead of `tablename_tp`, always `tablename` must be referenced.
5. A View may reference other views, provided that referenced views occur *earlier* in `$SCHEMA`. This way, cyclic forward-references are impossible.

A view is constructed by adding the following to `$SCHEMA`:

`"viewname" => array("VIEW" => query)` The schema compiler will automatically determine the corresponding `FIELDS` and other descriptions as you would normally supply in a table definition. You cannot provide other definitions by hand; they will be ignored. The *query* syntax is described in sections 5.3.4 and 5.3.5.

2.2 \$CONFIG configuration

The variable `$CONFIG` in `config.php` describes the databases of a system and the connection methods for working with them. It is a hash with the following components:

`USE_AUTH => true` or `false`. Defaults to `true` when omitted.

When set, the whole system works with profile-based user authentication and authorization. The following tables are automatically added to `$SCHEMA`:

profiles A temporal table describing the access rights for a class of users. This table contains the following columns:

`profile_name` Secondary key, uniquely describing each profile.

`profile_descr` Description, for comments and remarks.

`t_tablename` columns of this form are automatically added and maintained by `create_schema.php` for each other table in the system. It specifies the access rights to the respective table in the same format as described in section 2.1.1. The difference to the specification in `$SCHEMA` is that here we define *dynamic* access rights for each profile, which may *lower* the rights specified in the *static* `$SCHEMA`. This way, some groups of users may be restricted to readonly access, or be even denied access to whole tables.

`f_tablename_fieldname` Automatically added and maintained by `create_schema.php` for all columns of all other tables in the system. It specifies the access rights at column granularity. It may further restrict the access rights as already restricted by `t_tablename`.

languages A temporal table which may be used for internationalization of the user interface. Actually, it just defines which templates to use, so it may be also used for creation of different *skins* (look-and-feel styles) in the same language. It contains the following columns:

language_name The secondary key uniquely describing the interface style.

language_template The name of the compiled template file to use. Defaults to `generic_generic.php`. Details on the template compiler are in section 3.3.

language_descr Description, for comments and remarks.

users A temporal table having the following columns:

user_name Secondary key, uniquely describing each user.

user_password This is stored in encrypted form (see the `password()` function in MySQL).

profile_name References **profiles**. Thus it determines the access rights for the user.

language_name References **languages**. Thus it determines the user interface for the user.

user_descr Description, for comments and remarks.

When **USE_AUTH** is **false**, these tables are neither created nor used. The login authentication process just passes the username and password credentials to the database and hopes to get a connection with the necessary access rights. Since there is no profile info about runtime access permissions, only the **ACCESS** attributes can be used for *static* customization of the UI.

USE_ENGINE => **true** or **false**. Defaults to **false** when omitted.

Specifies that additional tables for the business process engine shall be added to **\$SCHEMA**. Details are in section 4.

CONNECTIONS => *hash*. No default. Cannot be omitted.

Describes the connections to databases. It has the form *connection_name* => *hash*, where *connection_name* is an arbitrary internal identifier which may be used in DB specifications of **\$SCHEMA** (see section 2.1.1). The *hash* is two or more associations from the following:

MASTER => *string*. No default. Cannot be omitted.

Hostname of the database server.

DRIVER => *string*. Defaults to **“mysql”** when omitted.

Currently this is the only driver, but others are planned.

BASE => *string*. No default. Cannot be omitted.

Name of the database (see **show databases** command in most SQL dialects).

USER => *string*. No default. Can and should be omitted.

PASSWD => *string*. No default. Can and should be omitted.

Although you can provide usernames and passwords here, this is considered *bad practice*. When someone hacks your **wwwrun** account, he can read your passwords in plaintext. Instead, when you omit these parameters here, the login credentials from the Apache authentication is used as database access credentials at runtime. Many system designers try to avoid the overhead of maintaining database access rights for individual users. The Athomaris PHP Library *automates* this task: Whenever an entry in **users** is created or updated via `db_insert()` or `db_update()`, the database access rights (e.g. in MySQL this is the global table `mysql.user`) are automatically updated *in sync* with that. This works even when a **profile** is altered or when the foreign key **profile_name** is changed.

SLAVES => *array* of strings. No default. Can be omitted.

When using master/slave replication in MySQL, you can provide a list of slave servers. *Readonly* accesses are directed to (randomly selected) servers from the list, in preference to the master server. This way, the overall system load can be spread more smoothly. This is important in environments dominated by database reads. This feature is crucial for enterprise-grade scalability.

3 Basic Customization

3.1 Link-Headers

By default, each page starts with links to all tables from `$SCHEMA`. You can customize these and add your own links as follows:

Define a global variable `$LINKS` (best practice: do this in `config.php`) which is a hash indexed by categories. Each category will later be displayed in a different line. The default links are all in the default category `"Tables:"`. In turn, a link category is another hash indexed by *link names* (see `{textlist}` in section 3.3). In the default links, the link name is equal to the table name; in general you may create your own names as you like. Each link name may be associated with an arbitrary HTML href. For example, the default links for project `demo_basic` are constructed as follows:

- `http://localhost/demo_basic/index.php?table=foos`

Hint: if you want to remove some of the default links, just define something like

- `$LINKS["Tables:"]['foos'] = "";`

Of course, by defining a non-empty string you can redirect it to something else. The whole default category `Tables:` can be removed by

- `$LINKS["Tables:"] = "";`

Notice that `undef($LINKS["Tables:"])` will not work because it is equivalent to non-defining `$LINKS` at all, which triggers the default link creation. When you remove the default category, you should provide *all* links by hand, otherwise your application would be completely link-less.

3.2 Generic Application Layer

After inclusion of `common/app.php`, some high-level functions are available for easy PHP programming:

`app_get_templates()` This should be called once after you have included the application layer by `require_once("$BASDIR/./common/app.php")`. This loads the templates as configured for the currently logged-in user.

`app_links()` When called, the default links as defined by the global structure `$LINKS` (see section 3.1) are displayed.

`app_display_table("tablename")` Displays the table in browsable form. When the access permissions allow writing, buttons for manipulation of the data are also displayed.

`app_input_table("tablename")` Displays all fields for which the current user has write access, and handles all the corresponding input actions.

`app_display_record($table, $cond)` Displays a single record. The selection is specified by `$cond` as described in section 5.3.3. Upon write permissions, an input table is also added.

`app_display_download($table, $cond, $download_fieldname, $filename)` Starts the download of the contents of an arbitrary field (most useful for BLOBs). This must be called as the only function in a single request, because it generates an HTML header suggesting a filename for saving the data at the client side.

The following request parameters are used by the above functions:

`$_REQUEST["table"]` The table name we are currently working on.

`$_REQUEST["order"]` Only for `app_display_table()`. Specifies a comma-separated list of field names for ordering of the table.

`$_REQUEST["primary"]` Tells which fields are identifying the tuple for `app_display_record()` or `app_display_download()`. This must be comma-separated list of field names. You may specify *any* fields which can *together* uniquely identify the tuple - it need not be the primary key, but can also be an `UNIQUE` key. In addition, you also have to supply field values for these fields. For example, if you specified `$_REQUEST["table"]="foos"` and `$_REQUEST["primary"]="foo_id"`, you also have to supply `$_REQUEST["foo_id"]` in order to uniquely identify the tuple you want to work with.

`$_REQUEST["delete"]` When set, this is a delete request. The tuple is *always* identified by its primary key. Values for the primary key must also be submitted as additional `$_REQUEST` fields. For example, if table `foos` has primary key `foo_id`, the value of `$_REQUEST["foo_id"]` must be supplied to uniquely identify the record which should be deleted.

`$_REQUEST["change"]` When set, this is a change request for an existing tuple. In addition to values for the obligatory primary key, you can specify as many additional fields as you like to change. When you *omit* a field, its value is left unchanged (see also `db_update()` in section 5.3.2).

`$_REQUEST["insert"]` Like `change`, but a new tuple is always inserted (see also `db_insert()` in section 5.3.2). You should provide the values of all fields; when you omit one, the `DEFAULT` will be inserted instead. Supplying values for the primary key is useless because it is always initialized by `auto_increment`.

`$_REQUEST["edit"]` Tells `app_input_table()` which tuple it should prepare for editing. This does not yet alter the tuple, but just generate a form for altering by the user.

`$_REQUEST["clone"]` Like `edit`, but instead of altering an existing tuple, a new one is presented for editing, inheriting all field values from the original tuple.

`<none>` When neither `edit` nor `clone` is set, a new tuple will be presented for editing with `DEFAULT` input fields.

3.3 Generic Template Mechanism

Templates are used to separate the presentation layer from the business logic. In other frameworks, this is called a “Model-View-Controller” (MVC) architecture. While our model is maintained by `$$SCHEMA / $EXTRA` and the controller is simply the Athomaris PHP Library invoked by Apache, views (in that terminology) are implemented via our templates.

Templates reside in ASCII files having the extension `.tpl`. The template compiler (see section 1.1) finds all files with that extension and translates them into fast PHP code.

At the outermost level, a template file can contain the following elements:

`{include "filename" /}` Includes another file as if its contents were written at the insertion point. After that, you may *redefine* any templates or textlists. Since redefinitions replace the previous definition, you can easily create variants of template sets this way. In particular, new *skins* can be created rather quickly.

`{textlist}...{/textlist}` The content at ... must be a sequence of ASCII lines, each terminated by a newline. Each line defines a simple macro via the following syntax:

`name = substitution_text` The usage of simple macros is explained later (see element `{text "name"}`). Example: change the default link-header category `Tables:` (see section 3.1) to something else by

- `{textlist}`
- `Tables: = Hey, here you can browse the following tables:`
- `{/textlist}`

`{template "name"}...{/template}` Defines a new template. The template compiler translates this into a PHP function `tpl_name($data) {...}` which can be called like any ordinary PHP function. When called, this function simply outputs the text between `{template "name"}` and `{/template}`. The text may contain further directives and macro expansions (working on the parameter `$data`) as follows:

`{$var/}` Output the PHP parameter `$data["var"]`, as provided to the template invocation. If the substitution text contains any HTML special characters, they are quoted to avoid XSS attacks onto your system. Deeper levels of the `$data` structure may be accessed via the syntax `{$var->subvar}` which outputs `$data["var"] ["subvar"]`. Note that `{$var->$subvar}` would result in something different: `$data["var"] [$data["subvar"]]`. As a simple rule of thumb, just remember that the number of `$` signs in your code indicates nothing more but the number of substitutions carried out at runtime.

`{text "name"/}` Output the simple macro defined in a `{textlist}`. If `name` does not exist anywhere, its name is outputted unmodified in place of the non-existing definition (fallback). Note there is no quote-protection against HTML code, because the `{textlist}` should have been written by a trusted programmer who probably even *needs* to include raw HTML code for his purposes.

`{text $var/}` Same as before, but the `name` is (indirectly) fetched from `$data["var"]`. The same principle applies orthogonally to all examples where a string in double quotes is used: instead of the string constant, you may always use a variable instead, or vice versa.

`{ascii $var/}` Like `{$var}`, but the text is written in typewriter font (monospaced), blanks are translated to ` `, and newlines are translated into `
`. This way, a kind of "verbatim output" is produced. Handy for displaying raw computer data which should be protected from HTML interpretation.

`{preview $var/}` Same as `{ascii $var}`, but when the text is very long, it is abbreviated with `...` and a link is displayed. When the user clicks on the link, the full text will be displayed in a new window.

`{param $var/}` Like `{$var}`, but URL escaping is used instead of HTML escaping. Especially useful for generating URL parameters like `?paramname={param $myvalue/}`.

`{raw $var/}` Dangerous! Like `{$var}`, but bypasses the HTML quote protection. Use only if you really know what you are doing!

`{header $var/}` Generate a raw HTML header. This must be called before any other output (see also PHP function `header()`).

`{row $var/}` Only for very special cases. Assumes that `$data["var"]` is a structured data record. This outputs an escaped string which encodes the whole record, such that it can be decoded again via the library function `_tpl_decode_row()`.

`{printf "format" $var1 $var2 .../}` Obvious semantics for C programmers (see also the PHP function `sprintf()`). The output is quote-protected.

`{var $name = expression /}` Assign a new value to `$data["name"]`. The *expression* may be an arbitrary PHP expression which *must not contain any spaces*, and you can access only the variable `$data` and subfields thereof.

`{unset $name /}` Deletes the variable `$data["name"]`.

`{tpl "othername"/}` Calls the template `tpl_othername` with unchanged argument `$data`. The template `othername` must exist, otherwise an error is produced. If you want to submit a different argument structure to the callee, you can use one of the following variants:

`{tpl "name" ($otherdata)/}` Uses `$data["otherdata"]` instead of `$data`.

`{tpl "name" "key1" => "value1", "key2" => $var2, .../}` The callee is provided with new or altered values for the given keys. When you combine this with `{if}`, you can even do recursive calls.

`{hook "othername"}/` Like `{tpl}/`, but no error is thrown if the callee does not exist. This is handy for introducing self-documenting hooks, simplifying plugin architectures.

`{if condition1}...{elseif condition2}...{else}...{/if}` Obvious semantics for PHP programmers. Of course, the `{elseif}` and `{else}` parts may be omitted.

`{loop $var as $value}...{/loop}` The loop body is repeated for each substructure member of `$data["var"]`, and `$data["value"]` is assigned the corresponding value during the loop. After the loop has finished, `$data["value"]` is restored to its previous value.

`{loop $var as $key => $value}...{/loop}` Variant thereof: both `$data["key"]` and `$data["value"]` are assigned the hash key and the hash value of `$data["var"]` during the loop, respectively. Both are restored afterwards.

`{loop as $value}...{/loop}` Variant: by omitting `$var`, you can iterate over the top-level `$data` instead of over one of its substructures. Recommendation: try to avoid this in your designs, because later addition of sibling fields is impossible. Just use an additional dummy level like `$DATA`, even if you currently don't need it.

3.4 Default Templates

The template mechanism as described in the previous section may be used in almost arbitrary ways. However, application programming can greatly improve efficiency if some appropriate **conventions** are introduced.

The following conventions are used in `../common/app.php` and are recommended as a basis for your own extensions:

...

4 Business Process Engine

4.1 Concepts

Most business process languages such as BPEL follow a *procedural* paradigm for describing a workflow, similar to classical procedural programming languages. These models are extremely complex: the current WS-BPEL specification 2.0 prints to 140 pages DINA4 in very small font. In contrast, the Athomaris engine tries to remain simple, but to reach at least the same *expressive power*. In future releases, we want to implement translators from BPEL to our model.

The Athomaris business process engine uses the well-known concept of **finite automaton**s as established in computer science for decades. The most important difference to the theoretical concept is **explicit state**: *automatic instantiation* of new automaton, as soon as a new tuple is added to the database. In other words: a new finite automaton is instantiated automatically for *any* tuple of the database (whether it be a newly inserted tuple, or a previously modified tuple), working on that tuple independently from other automaton. The only triggering condition is that some *rule* must match on that tuple. Thanks to implicitly matching *any* existing tuple, the system can achieve almost arbitrary parallelism and scalability. If you want a purely sequential workflow, just don't produce *new* tuples which would be “watched” by the Athomaris engine - instead just update the explicit state which is kept inside your *old* tuple. Another benefit of explicit state is that other applications can query and work with it.

The Athomaris model is much more similar to a rule-based computation paradigm called “black-board systems”. Another similar model is called “coordination models and languages” which has attracted a lot of research some years ago.

Here is a small comparison of concepts between procedural models like BPEL and our non-procedural “descriptive” approach:

model	execution state	parallelism	waiting	control flow	data flow
procedural	implicit	explicit	explicit	explicit	implicit
rule-based	explicit	implicit	implicit	implicit	implicit

Note: this description is for technicians only. We deserve to write another (friendlier) description for business managers, employing their terminology and buzzwords.

4.2 Usage

The business process engine is enabled by `USE_BUSINESS_ENGINE => true` in `$CONFIG` (see section 2.2). This adds some orchestration tables to `$SCHEMA` as described in the next section. After you have filled those tables which orchestration rules or have updated them, you have to invoke the orchestration compiler via the following URL:

- `http://localhost/demo_business/orchestrator.php`

In a future release of Athomaris, we want to have a graphical tool under this URL where you can view and create orchestrations, probably even with drag and drop. For now, there is only a batchmode compiler.

When compilation is successful, a file `compiled/engine_table.php` is created. Afterwards, you may invoke the business process engine from the commandline by supplying username and password as arguments:

- `cd /www/athomaris/demo_business; php business_engine.php “root” “secret”`

4.3 Tables for the Orchestration Level

Once the business process engine has been enabled via the `USE_BUSINESS_ENGINE` attribute of `$CONFIG` (see section 2.2), the following tables are automatically added to the schema:

bps (abbreviation for *business processes*). This table describes the *interface* to a business process. When used standalone without joining to other tables such as **rules**, you get an *abstract* business process (borrowing terminology from other languages such as BPEL). However, when you join this with **rules** and **conts** (see below), you get an *executable* business process. In addition to the usual temporal columns such as `bp_id` and `bp_version`, this table consists of the following columns:

bp_name Secondary key. This is used to uniquely identify each business process in the system.

bp_statefield Defines on which *cell candidates* the business process should work. Must be a string of the form *tablename.fieldname* where *tablename* is a valid name of another table from `$$SCHEMA` and *fieldname* a valid column name therein.

Currently not yet implemented: if you leave this field empty, a kind of “procedural model” will be executed. State is chosen automatically, but the business process needs to be *called explicitly*.

bp_inputs, **bp_outputs** Not yet implemented. In future releases, this will be used for argument checking when calling procedural-style business processes and sub-processes. Leave it empty for now.

bp_joinwith Usually left empty. When you specify a comma-separated list of other table names, these tables will be naturally joined with *tablename*. In consequence, you may use their fields in `@{otherfield}` macros (see below). This way you can easily access related data.

bp_comment For documentation at the orchestration level.

rules This table describes the left-hand part of an automaton rule. In addition to the usual temporal columns such as `rule_id` and `rule_version`, it consists of the following columns:

bp_name References **bps**. Tells to which business process the current rule belongs.

rule_prio This is used to define an *order* on all rules working for the same business process. In other words, the combined secondary key of **rules** is defined as `UNIQUE => array('bp_name', rule_prio')`.

rule_startvalue Defines which value a *cell candidate* must have if this rule shall “fire”. Once the rule has fired on a candidate, the execution engine remembers the identity of the firing cell (in the following, we mean this single instance when we speak of “cell”). In order to allow multiple modes of testing and matching against cell candidates, the following syntax must be used:

`=value` Tests the candidate for equality with the given constant.

`%value` Matches if *value* is a substring of the candidate cell.

`/regex/` Matches *regex* against the candidate cell. You can use Perl-compatible regular expressions with parentheses. Later these may be used for `@{n}` macro substitutions (see **rule_action** below). Note that full-line matches must be indicated by `/\A...\Z/` or similar.

rule_condition Defines additional conditions which must hold if the rule shall fire. The syntax is as follows:

<empty> When left empty, the condition is always true, i.e. there is no additional condition.

?fieldname=value Similar to **rule_startvalue**, but you may test arbitrary other fields. If you have used natural joins at **bp_joinwith**, you can use them here. This is very useful for formulating complex **rule_conditions**. You may denote the same variants of tests as with **rule_fieldvalue**, for example **?fieldname%value** or **?fieldname/regex/**.

You may specify multiple conditions on different lines; these are logically anded together. In a future release, we plan to support full-fledged boolean expressions.

rule_location Not yet implemented. This will specify the hostname or a group of hostnames where the rule should be executed.

rule_firevalue As soon as the rule fires, this value will be immediately written to the cell. This is used for proper restarting in case of system crashes. It simply records the fact that a rule has fired and thus an action has started (and probably must be rolled back). You may provide one of the following syntaxes:

- <empty>** When left empty, a reasonable default is written into the cell. If the cell contained an integer number, this number is incremented by 1. If it contained a string, the string is *prepended* with “start_”. Recommendation: don’t use the default, always specify this value explicitly!
- + *n*** When starting with a plus sign followed by a blank and a number, this number is added to the old cell contents. Works only well if the cell had an integer type.
- <anything else>** Will be literally copied to the cell. However, it may contain @ macros, which are *substituted* as follows:
 - @{*fieldname*}** This is substituted by the value of *fieldname*. It may stem from the matching tuple, or even from a more complex join delivered by **bp_joinwith**.
 - @{*n*}** When *n* is a number, it is substituted by the *n*-th matching parenthesis of the regular expression denoted in **rule_startvalue**. This way, you can build your own parser and propagate almost arbitrary values from there.
 - @(*subcommand args...*)** This is substituted with the standard output of *subcommand*, executed in a subshell. For example, you may do simple calculations such as **@(expr @{foo_state} + 1)**.

rule_action Tells what to do when the rule fires. The following kinds of actions are currently implemented (further actions such as SOAP are planned):

- script *command args...*** Executes the given command with parameters in a Unix shell. Inside your text, you may use the same macro substitutions as explained earlier.
- url *http://...*** This is equivalent to **script wget -O - 'http://...'**. You can query arbitrary web or ftp servers this way. When combined with the macro features explained above, you may easily create arbitrary REST queries, just for example.
- insert *othertablename fieldname1='value1', fieldname2='@{something}'...***
Inserts a new tuple into the database. Notice: when some (other) rules are defined for this table, this effectively results in creating a new non-procedural sub-business-process.
- update *fieldname1='value1', fieldname2='@{something}'...*** Update the original candidate tuple the current rule is working on.
- update *tablename fieldname1='value1', fieldname2='@{something}'...*** Same, but update a tuple from *tablename*. When *tablename* is mentioned in **bp_joinwith**, you need not supply the primary key because it can be determined automatically. Otherwise, the primary key of the tuple must be explicitly given.
- delete *tablename fieldname1='value1'...*** As the name suggests, a tuple is deleted from the database. The primary key must be always given. Be careful with it!

Multiple actions may be supplied if they are on different lines. They will be executed sequentially. Note that a failing action (such as a database error) will immediately terminate the action chain - the following actions will not be started in such a case.

rule_timeout When non-zero, each started business process is monitored. Whenever it does not respond either on **stdout** or **stderr** within the given timeout (measured in seconds), a pseudo-event **TIMEOUT** is generated (see section 4.4). The pseudo-event can then be handled as described below.

rule_comment For documentation at the orchestration level.

Note that table **rules** describes only which business process to start under some preconditions. The *consequences* of a business process execution are solely handled by *continuations*:

conts (abbreviation for *continuations*). This table describes the *possible consequences* of a rule execution. In addition to the usual temporal columns such as **cont_id**, it consists of the following columns:

- bp_name, rule_prio** Together they references **rules**. Tells to which rule the current continuation belongs.
- cont_prio** This is used to define an order on all continuations which are *candidates* for handling the same rule invocation. **conts** has the following secondary key: **UNIQUE => array("bp_name,rule_prio,cont_prio")**.

cont_match Regular expression which matches *each line* from **stdout** and **stderr** of the **rule_action**, as well as pseudo events such as **TIMEOUT *n*** or **STATUS *n*** (see section 4.4). When the regular expression matches, the continuation is *selected*.

cont_action Same syntax as described at **rule_action**. Although it is possible to start another scripts again, their output is not parsed again. So use this only for small actions not involving different state changes at the cell.

cont_endvalue After this cont has been selected and the **cont_action** has successfully completed, this value is written back into the cell. Used to record the fact that the action has completed. When you leave this empty, no value is written and the next continuation according to **cont_prio** is examined whether this can be selected (continue the candidate selection process). When non-empty, you may use the “+ *n*” notation and macros as described at **rule_firevalue**.

cont_comment For documentation at the orchestration level.

4.4 Pseudo Events

A pseudo event is generated by the business process execution engine, and it is treated *uniformly* in the same way as a line of output from the script will be treated. Currently the following events are defined:

START *pid* Tells that a script has started with process id *pid*.

STATUS *n* Tells that the script has terminated with Unix status **exit(*n*)**. Afterwards, you will never receive any further events.

INSERT *ok* Tells whether an **insert** action was successful.

UPDATE *ok* Tells whether an **update** action was successful.

DELETE *ok* Tells whether a **delete** action was successful.

TIMEOUT *n* Tells that the action is hanging or the script has not responded in any way for at least *n* seconds.

5 Advanced Features

5.1 Automatic Data Synchronization

...

5.2 Subrecords and Display/Editing of References

...

5.3 Programmers API to the Database

This section is devoted to experienced PHP programmers for whom the default `app.php` is not enough, or if they like to program their own interfaces or complex database access engines.

The purpose of the following functions is abstracted access to temporal tables, which is independent from SQL dialects.

5.3.1 Data Format

All database operations use a common PHP representation for table data, called `$data`. It is simply an *array of records*. A record is simply a hash, indexed by the column names. For example:

```
$data =
    array(
        array(
            'foo_id' => 17,
            'foo_comment' => 'first_example_record',
        ),
        array(
            'foo_id' => 18,
            'foo_comment' => 'next_example_record',
        ),
    );
```

As you can see, the column names are *deliberately* repeated. Although OO classes *could* save *some* of that runtime space overhead *in theory*, maintaining OO class definitions instead would be a greater effort, and it would lower flexibility. In particular, `ACCESS => 'n'` (see section 2.1.1) will omit inaccessible fields *dynamically at runtime*, individually for each field. Another use is for `db_update()`, see section 5.3.2: There you can simply specify *which* fields to update, just by omitting all those fields you don't want to change. Note that in absence of database transactions, this has different *atomic properties* from first reading a full tuple, updating the PHP structure, and finally writing back the full tuple. The difference becomes clear with MySQL MyISAM tables in presence of *concurrent* updating.

This kind of flexibility is hardly to achieve with the usually static OO class concepts. Thus we use a “functional” programming style (inspired by languages like ML, Haskell or good old Lisp), and we pay some (acceptable) runtime space overhead in favor of a simple design. Nevertheless, future releases of the Athomaris PHP library could support automatic generation of class definitions from `$SCHEMA`.

5.3.2 Update Operations

`$ok = db_insert('table', $data);` This will insert all the records from `$data` (see section 5.3.1) into the named table. Atomicity is only guaranteed for the consistency of each tuple; however

depending on the database type and the network distribution, atomicity for the whole set of records could be nevertheless achieved or at least approximated.

`$ok = db_update("table", $data);` This will update all *present* fields of all records from `$data`, with one exception: all fields from the PRIMARY key *must* be present and are *not* updated; otherwise the tuples could not be uniquely identified.

`$ok = db_replace("table", $data);` When the tuples are not yet present in the databases, this will lead to the same effect as `db_insert()`. Otherwise, the effect of `db_update()` will be achieved.

5.3.3 Reading of Data

`$data = db_read("table", "field1,field2,...", $cond, $order, $start, $count);`

This will retrieve data from the named table. The parameters have the following meaning:

`"field1,field2,..."` When this parameter is empty, all *accessible* fields (at least ACCESS => "r", see section 2.1.1) will be retrieved.

`$cond` When empty or null, all tuples from the whole database are retrieved. In its simplest form, you may provide a hash `array("fieldname" => fieldvalue)` with the obvious meaning: only tuples exactly matching the *fieldvalue* are selected. When providing multiple hash keys, all conditions are logically anded. More complex boolean expressions may be constructed in the following way:

1. the variant `"fieldname operator" => fieldvalue` uses *operator* instead of the default equality. Following *operators* are possible:

`= <> < > <= >=` The usual comparison operators.

`%` Tests for *like* as defined by ANSI SQL.

`rlike` Tests for MySQL *rlike*. Not portable to other databases! Try to avoid this.

`!` Tests whether *fieldname* is SQL NULL. The *fieldvalue* is ignored.

`@` Tests whether *fieldname* is *not* SQL NULL. The *fieldvalue* is ignored.

`in` The *fieldvalue* must be a sub-query as described in sections 5.3.4 or 5.3.5. The meaning is intuitively clear to SQL programmers.

`exists` The *fieldname* must be empty, and the *fieldvalue* must be a sub-query as described in sections 5.3.4 or 5.3.5. The meaning is intuitively clear to SQL programmers.

`not exists` Similar; the meaning is intuitively clear to SQL programmers.

1. when *fieldvalue* is an *aggregating sub-query* as described in section 5.3.5, a nested natural join with the *table* is computed (resulting in a *dependent* subquery), and its result is compared with *fieldname*. An operator may be appended to *fieldname*.
2. when `"fieldname"=>` is omitted and *fieldvalue* is itself a hash obeying the same construction rules, these conditions are logically *ored* instead of anded. When recursively nesting such expressions, oring and anding will always *alternate*, depending on odd or equal nesting level.

`$order` When empty, the result tuple may be unordered. Otherwise it must be a *string*, containing a comma-separated list of fieldnames. Alternatively, it may be an array of plain fieldnames.

`$start` When empty or 0, the result set is starting at its beginning. Otherwise it must be a number indicating the start position.

`$count` When empty or 0, the number of tuples is not restricted. Otherwise at most `$count` tuples will be delivered.

5.3.4 Full Queries / Subqueries

Queries and subqueries may be uniformly expressed as a single *hash* containing the following keys:

TABLE => *“table”*. Specifies the table to query. When a comma-separated list of table names is provided, the *natural join* of all tables is computed.

FIELD => *\$struct*. When empty, all fields are retrieved. When *\$struct* is a plain string containing a comma-separated list of fieldnames, only those fields are retrieved. When it is a hash, the following variants are possible:

“fieldname” Retrieve that field.

“aliasname” => *\$subquery* The subquery must be recursively structured in the same way as described in this section. The inner query is automatically *dependent* from the outer one by all common fieldnames (similar to a *natural join*), and the result of the subquery is propagated to the outer query under the name *aliasname*. This way, you may create “virtual fields” containing *dynamically computed* values.

COND => *\$cond*. Query condition, as already described in section 5.3.3. Recursive nesting is possible.

ORDER => *\$order*. Already described in section 5.3.3.

START => *\$start*. Already described in section 5.3.3.

COUNT => *\$count*. Already described in section 5.3.3.

JOINFIELDS => *“fieldname1,fieldname2,...”*. May be omitted. Defaults to *all* fieldnames occurring in more than one table from the **TABLE** list.

Use these fields for constructing a natural join. Each of the fieldnames must occur in at least two of the tables mentioned by **TABLE**.

JOIN_ON => *array* of strings. May be omitted. Defaults to *all* possible conditions derivable from **JOINFIELDS**, expressed in the following syntax:

“tablename1.fieldname1=tablename2.fieldname2” The tablenames must occur in **TABLE**, and each fieldname must occur in the corresponding table. When *tablename1* and *tablename2* are different, non-natural joins may be expressed. When you provide **JOIN_ON**, **JOINFIELDS** is ignored as a whole. By supplying an empty array, you may express the full cartesian product between the tables.

Instead of calling `db_read()`, you may use `_db_read()` the following way:

```
$data = db_read("table", "field1,field2,...", $cond, $order, $start,
$count); is equivalent to $data = _db_read(array("TABLE" => "table", "FIELD"
=> "field1,field2,...", "COND" => $cond, "ORDER" => $order, "START" => $start,
"COUNT" => $count);
```

5.3.5 Aggregated Queries / Subqueries

The keys **TABLE**, **COND**, **ORDER**, **START** and **COUNT** have been already described in the previous section. Instead of **FIELD**, you provide the following key:

AGG => *\$struct*. This is another *hash* consisting of the following two members:

FIELD => *“expression”*. Allowable is `min(*)`, `max(*)`, `count(*)` and `avg(*)`.

GROUP => *“fieldlist”*. This specifies the fields over which aggregation will be performed. The meaning is intuitively clear to SQL programmers knowing **group by** statements.

In addition, **JOINFIELDS** or **JOIN_ON** may be added as described in section 5.3.4.